

# INT305 note

(Machine Learning)


## 1 Introduction

### 1.1 Supervised learning (much of this course)

Task	Inputs	Labels
object recognition	image	object category
image captioning	image	caption
document classification	text	document category
speech-to-text	audio waveform	text
$\vdots$	$\vdots$	$\vdots$

#### 1.1.1 KNN

- Nearest neighbours **sensitive to noise or mis-labelled data** ("class noise").
- Smooth by having  $k$  nearest neighbours vote.

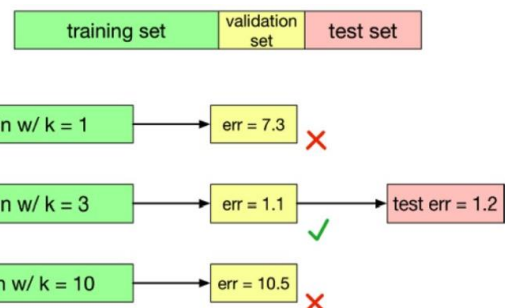



Algorithm (kNN):

1. Find  $k$  examples  $\{\mathbf{x}^{(i)}, t^{(i)}\}$  closest to the test instance  $\mathbf{x}$
2. Classification output is majority class

$$y = \arg \max_{t^{(z)}} \sum_{i=1}^k \mathbb{I}(t^{(z)} = t^{(i)})$$

- Balancing hyperparameter  $k$ 
  - Optimal choice of  $k$  depends on number of data points  $n$ .
  - Nice theoretical properties if  $k \rightarrow \infty$  and  $k/n \rightarrow 0$ .
  - Rule of thumb: choose  $k < \sqrt{n}$ .
  - We can choose  $k$  using validation set.



## 2 Linear Methods for Regression, Optimization

Linear regression exemplifies recurring themes of this course:

- Choose a **model** and a **loss function**
- Formulate an **optimization problem**
- Solve the minimization problem using one of two strategies

➤ **Direct solution** (set derivatives to zero)

➤ **Gradient descent**

- **Vectorize** the algorithm, i.e. represents in terms of linear algebra
- Make a linear model more powerful using features
- Improve the generalization by adding a **regularizer**

## 2.1 Supervised Learning Setup

In supervised learning:

- There is input  $\mathbf{x} \in \mathcal{X}$ , typically a vector of features (or covariates)
- There is target  $t \in \mathcal{T}$ , (also called response, outcome, output, class)
- Objective is to learn a function  $f: \mathcal{X} \rightarrow \mathcal{T}$  such that  $t \approx y = f(\mathbf{x})$  based on some data  $\mathcal{D} = \{(\mathbf{x}^{(i)}, t^{(i)}) \text{ for } i = 1, 2, \dots, N\}$


## 2.2 Linear Regression

### 2.2.1 Linear Regression Model

**Model:** In linear regression, we use a linear function of the features  $\mathbf{x} = x_1, \dots, x_D \in \mathbb{R}^D$  to make predictions  $y$  of the target value  $t \in \mathbb{R}$ :

$$y = f(\mathbf{x}) = \sum_j w_j x_j + b$$

- $y$  is the **prediction**
- $\mathbf{w}$  is the **weights**
- $b$  is the **bias** (or intercept)
- $\mathbf{w}$  and  $b$  together are the **parameters**
- We hope that our prediction is close to the target:  $y \approx t$ .



- If we have only 1 feature:  $y = wx + b$  where  $w, x, b \in \mathbb{R}$ .
- $y$  is linear in  $x$ .
- If we have only  $D$  features:  $y = \mathbf{w}^T \mathbf{x} + b$  where  $\mathbf{w}, \mathbf{x} \in \mathbb{R}^D, b \in \mathbb{R}$
- $y$  is linear in  $\mathbf{x}$ .

### 2.2.2 Linear Regression workflow

We have a dataset  $\mathcal{D} = \{(\mathbf{x}^{(i)}, t^{(i)}) \text{ for } i = 1, 2, \dots, N\}$ :

- $\mathbf{x}^{(i)} = (x_1^{(i)}, x_2^{(i)}, \dots, x_D^{(i)})^T \in \mathbb{R}^D$  are the inputs (e.g. age, height)
- $t^{(i)} \in \mathbb{R}$  is the target or response (e.g. income)
- Predict  $t^{(i)}$  with a linear function of  $\mathbf{x}^{(i)}$ :
  - $t^{(i)} \approx y^{(i)} = \mathbf{w}^T \mathbf{x}^{(i)} + b$
  - Different  $\mathbf{w}, b$  define different lines.
  - We want the "best" line  $\mathbf{w}, b$ .

### 2.2.3 Linear Regression Loss Function

• A loss function  $\mathcal{L}(y, t)$  defines how bad it is if, for some example  $\mathbf{x}$ , the algorithm predicts  $y$ , but the target is actually  $t$ .

• **Squared error loss function:**

$$\mathcal{L}(y, t) = \frac{1}{2} (y - t)^2$$

- $y - t$  is the residual, and we want to make this small in magnitude.
- The  $1/2$  factor is just to make the calculations convenient.

- **Cost function:** loss function averaged over all training examples.

$$\mathcal{J}(\mathbf{w}, b) = \frac{1}{2N} \sum_{i=1}^N (y^{(i)} - t^{(i)})^2 = \frac{1}{2N} \sum_{i=1}^N (\mathbf{w}^T \mathbf{x}^{(i)} + b - t^{(i)})^2$$

### 2.2.3 Linear Regression Vectorization

But if we expand  $y^{(i)}$ , it will get messy:

$$\frac{1}{2N} \sum_{i=1}^N \left( \sum_{j=1}^D (w_j x_j^{(i)}) + b - t^{(i)} \right)^2$$

**Vectorize** algorithms by expressing them in terms of vectors and matrices:

$$\mathbf{w} = (w_1, \dots, w_D)^T \quad \mathbf{x} = (x_1, \dots, x_D)^T$$

$$y = \mathbf{w}^T \mathbf{x} + b$$

Python code:

```

y = b
for j in range(M):
    y += w[j] * x[j]
    
```

$$=$$

```

y = np.dot(w, x) + b
    
```

Organize all the training examples into a **design matrix X** with one row per training example, and all the targets into the **target vector t**:

$$\mathbf{X} = \begin{pmatrix} \mathbf{x}^{(1)T} \\ \mathbf{x}^{(2)T} \\ \mathbf{x}^{(3)T} \end{pmatrix} = \begin{pmatrix} 8 & 0 & 3 & 0 \\ 6 & -1 & 5 & 3 \\ 2 & 5 & -2 & 8 \end{pmatrix}$$

one feature across  
all training examples

one training  
example (vector)

Computing the **predictions** for the whole dataset:

$$\mathbf{X}\mathbf{w} + b\mathbf{1} = \begin{pmatrix} \mathbf{w}^T \mathbf{x}^{(1)} + b \\ \vdots \\ \mathbf{w}^T \mathbf{x}^{(N)} + b \end{pmatrix} = \begin{pmatrix} y^{(1)} \\ \vdots \\ y^{(N)} \end{pmatrix} = \mathbf{y}$$

Computing the **squared error cost** across the whole dataset:

$$\mathbf{y} = \mathbf{X}\mathbf{w} + b\mathbf{1}$$

$$\mathcal{J} = \frac{1}{2N} \|\mathbf{y} - \mathbf{t}\|^2$$

We can also add a column of 1's to design matrix, combine the bias and the weights, and conveniently write:

$$\mathbf{X} = \begin{bmatrix} 1 & [\mathbf{x}^{(1)}]^T \\ 1 & [\mathbf{x}^{(2)}]^T \\ \vdots & \vdots \end{bmatrix} \in \mathbb{R}^{N \times (D+1)} \quad \text{and} \quad \mathbf{w} = \begin{bmatrix} b \\ w_1 \\ w_2 \\ \vdots \end{bmatrix} \in \mathbb{R}^{D+1}$$

Then, our predictions reduce to  $\mathbf{y} = \mathbf{X}\mathbf{w}$ .

## 2.3 Direct Solution •

### 2.3.1 Linear Algebra •

- We seek  $\mathbf{w}$  to minimize  $\|\mathbf{X}\mathbf{w} - \mathbf{t}\|^2$ , or equivalently  $\|\mathbf{X}\mathbf{w} - \mathbf{t}\|$
- $\text{range } \mathbf{X} = \{\mathbf{X}\mathbf{w} | \mathbf{w} \in \mathbb{R}^D\}$  is a  $D$ -dimensional subspace of  $\mathbb{R}^N$
- Recall that the closest point  $\mathbf{y}^* = \mathbf{X}\mathbf{w}^*$  in subspace  $\text{range}(\mathbf{X})$  of  $\mathbb{R}^N$  to arbitrary point  $\mathbf{t} \in \mathbb{R}^N$  is found by orthogonal projection.
- We have  $(\mathbf{y}^* - \mathbf{t}) \perp \mathbf{X}\mathbf{w}, \forall \mathbf{w} \in \mathbb{R}^D$
- $\mathbf{y}^*$  is the closest point to  $\mathbf{t}$



### 2.3.2 Calculus •

- **Partial derivative**: derivative of a multivariate function with respect to one of its arguments.

$$\frac{\partial}{\partial x_1} f(x_1, x_2) = \lim_{h \rightarrow 0} \frac{f(x_1 + h, x_2) - f(x_1, x_2)}{h}$$

- To compute, take the single variable derivative, pretending the other arguments are constant.
- Example: partial derivatives of the prediction  $y$ .

$$\begin{aligned} \frac{\partial y}{\partial w_j} &= \frac{\partial}{\partial w_j} \left[ \sum_{j'} w_{j'} x_{j'} + b \right] & \frac{\partial y}{\partial b} &= \frac{\partial}{\partial b} \left[ \sum_{j'} w_{j'} x_{j'} + b \right] \\ &= x_j & &= 1 \end{aligned}$$

- For **loss derivatives**, apply the chain rule:

$$\begin{aligned} \frac{\partial \mathcal{L}}{\partial w_j} &= \frac{d\mathcal{L}}{dy} \frac{\partial y}{\partial w_j} & \frac{\partial \mathcal{L}}{\partial b} &= \frac{d\mathcal{L}}{dy} \frac{\partial y}{\partial b} \\ &= \frac{d}{dy} \left[ \frac{1}{2} (y - t)^2 \right] \cdot x_j & &= y - t \\ &= (y - t) x_j & & \end{aligned}$$

- For cost derivatives, use linearity and average over data points:

$$\frac{\partial \mathcal{J}}{\partial w_j} = \frac{1}{N} \sum_{i=1}^N (y^{(i)} - t^{(i)}) x_j^{(i)} \quad \frac{\partial \mathcal{J}}{\partial b} = \frac{1}{N} \sum_{i=1}^N y^{(i)} - t^{(i)}$$

- Minimum must occur at a point where partial derivative are zero.

$$\frac{\partial \mathcal{J}}{\partial w_j} = 0 \ (\forall j), \quad \frac{\partial \mathcal{J}}{\partial b} = 0.$$

(If  $\partial \mathcal{J} / \partial w_j \neq 0$ , you could reduce the cost by changing  $w_j$ )

- We call the vector of partial derivatives the **gradient**.
- Thus, the "gradient of  $f: \mathbb{R}^D \rightarrow \mathbb{R}$ ", denoted  $\nabla f(\mathbf{w})$ , is:

$$\left( \frac{\partial}{\partial w_1} f(\mathbf{w}), \dots, \frac{\partial}{\partial w_D} f(\mathbf{w}) \right)^\top$$

- The gradient points in the direction of the greatest rate of increase.
- Analogue of second derivative (the "Hessian" matrix):

$$\nabla^2 f(\mathbf{w}) \in \mathbb{R}^{D \times D} \text{ is a matrix with } [\nabla^2 f(\mathbf{w})]_{ij} = \frac{\partial^2}{\partial w_i \partial w_j} f(\mathbf{w})$$

- We seek  $\mathbf{w}$  to minimize  $\mathcal{J}(\mathbf{w}) = ||\mathbf{X}\mathbf{w} - \mathbf{t}||^2 / 2$
- Taking the gradient with respect to  $\mathbf{w}$  we get:

$$\nabla_{\mathbf{w}} \mathcal{J}(\mathbf{w}) = \mathbf{X}^\top \mathbf{X} \mathbf{w} - \mathbf{X}^\top \mathbf{t} = 0$$


$$\mathbf{w}^* = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{t}$$

- **Linear regression** is one of **only** a handful of models in this course that **permit direct solution**.

## 2.4 Polynomial Feature Mapping

### 2.4.1 Introduction

The relation between the input and output may not be linear. But we can still use linear regression by mapping the input features to another space using **feature mapping** (or **basis expansion**).  $\varphi(\mathbf{x}) : \mathbb{R}^D \rightarrow \mathbb{R}^d$  and treat the mapped feature (in  $\mathbb{R}^d$ ) as the input of a linear regression procedure.




Find the data using a degree-M polynomial function of the form:

$$y = w_0 + w_1x + w_2x^2 + \dots + w_Mx^M = \sum_{i=0}^M w_i x^i$$

- Here the feature mapping is  $\varphi(x) = [1, x, x^2, \dots, x^M]^T$ .
- We can use linear regression to find  $\mathbf{w}$ , since  $y = \varphi(x)^T \mathbf{w}$  is linear with  $w_0, w_1, \dots, w_M$

### 2.4.2 Model Complexity and Generalization

- Underfitting (M=0): model is too simple – does not fit the data.
- Overfitting (M=9): model is too complex – fits perfectly.



### 2.4.3 L<sup>2</sup> Regularization

**Regularizer:** a function that quantifies how much we prefer one hypothesis VS another.

- We can encourage the weights to be small by choosing as our regularizer the **L<sup>2</sup> penalty**.


$$\mathcal{R}(\mathbf{w}) = \frac{1}{2} \sum_j w_j^2$$

(To be precise, the L<sup>2</sup> norm is Euclidean distance, we're regularizing the squared L<sup>2</sup> norm)

- The **regularized cost function** makes a tradeoff between fit to the data and the norm of the weights.

$$J_{reg}(\mathbf{w}) = J(\mathbf{w}) + \lambda \mathcal{R}(\mathbf{w}) = J(\mathbf{w}) + \frac{\lambda}{2} \sum_j w_j^2$$

- If you fit training data poorly,  $J$  is large. If your optimal weights have high values,  $\mathcal{R}$  is large.
- Large  $\lambda$  penalize wight values more.
- Like  $M$ ,  $\lambda$  is a hyperparameter we can tune with a validation set.



### 2.4.4 L<sup>2</sup> Regularized Least Squares: Ridge regression

For the least squares problem, we have  $J(\mathbf{w}) = \frac{1}{2N} \|\mathbf{X}\mathbf{w} - \mathbf{t}\|^2$

- When  $\lambda > 0$  (with regularization), regularized cost gives:

$$\begin{aligned} \mathbf{w}_\lambda^{\text{Ridge}} &= \underset{\mathbf{w}}{\operatorname{argmin}} J_{reg}(\mathbf{w}) = \underset{\mathbf{w}}{\operatorname{argmin}} \frac{1}{2N} \|\mathbf{X}\mathbf{w} - \mathbf{t}\|_2^2 + \frac{\lambda}{2} \|\mathbf{w}\|_2^2 \\ &= (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I})^{-1} \mathbf{X}^T \mathbf{t} \end{aligned}$$


- The case  $\lambda = 0$  (no regularization) reduces to least squares solution!

## 2.5 Gradient Descent

### 2.5.1 Concepts

- Many times, we do not have a direct solution: Taking derivatives of  $J$  w.r.t  $\mathbf{w}$  and setting them to 0 doesn't have an explicit solution.

- Gradient descent is an iterative algorithm, which means we apply an update repeatedly until some criterion is met.
- We initialize the weights to something reasonable (e.g., all zeros) and repeatedly adjust them in the direction of steepest descent.



(就是等到斜率为 0 即为最优解)

- Observe:
  - If  $\partial J / \partial w_j > 0$ , then increasing  $w_j$  increases  $J$ .
  - If  $\partial J / \partial w_j < 0$ , then increasing  $w_j$  decreases  $J$ .
- The following update always decreases the cost function for small enough  $\alpha$  (unless  $\partial J / \partial w_j = 0$ ):
- $\alpha > 0$  is a learning rate (or step size). The larger it is, the faster  $\mathbf{w}$  changes (but values are typically small).
- This gets its name from the gradient:

$$\nabla_{\mathbf{w}} J = \frac{\partial J}{\partial \mathbf{w}} = \begin{pmatrix} \frac{\partial J}{\partial w_1} \\ \vdots \\ \frac{\partial J}{\partial w_D} \end{pmatrix}$$

- Update rule in **vector form**:

$$\mathbf{w} \leftarrow \mathbf{w} - \alpha \frac{\partial J}{\partial \mathbf{w}}$$

- And for **linear regression** we have:

$$\mathbf{w} \leftarrow \mathbf{w} - \frac{\alpha}{N} \sum_{i=1}^N (y^{(i)} - t^{(i)}) \mathbf{x}^{(i)}$$

- So gradient descent updates  $\mathbf{w}$  in the direction of fastest decrease.
- Observe that once it converges, we get a **critical point**. i.e.  $\frac{\partial J}{\partial \mathbf{w}} = \mathbf{0}$

### 2.5.2 Gradient Descent for Linear Regression

- Why gradient descent, if we can find the optimum directly?
  - gradient descent can be applied to a much broader set of models
  - gradient descent can be easier to implement than direct solutions
  - For regression in high-dimensional space, gradient descent is more efficient than direct solution

### 2.5.3 Gradient Descent under the $L^2$ Regularization


- The gradient descent update to minimize the  $L^2$  regularized cost  $J + \lambda \mathcal{R}$  results in **weight decay**:

$$\begin{aligned} \mathbf{w} &\leftarrow \mathbf{w} - \alpha \frac{\partial}{\partial \mathbf{w}} (J + \lambda \mathcal{R}) \\ &= \mathbf{w} - \alpha \left( \frac{\partial J}{\partial \mathbf{w}} + \lambda \frac{\partial \mathcal{R}}{\partial \mathbf{w}} \right) \\ &= \mathbf{w} - \alpha \left( \frac{\partial J}{\partial \mathbf{w}} + \lambda \mathbf{w} \right) \\ &= (1 - \alpha \lambda) \mathbf{w} - \alpha \frac{\partial J}{\partial \mathbf{w}} \end{aligned}$$

### 2.5.4 Learning Rate (Step Size)

- In gradient descent, the learning rate  $\alpha$  is a hyperparameter we need to tune.
- Good values are typically between 0.001 and 0.1.





- To diagnose optimization problems, it's useful to look at **training curves**: plot the training cost as a function of iteration.
- Warning: in general, it's very hard to tell from the training curves whether an optimizer has converged. They can reveal major problems, but they can't guarantee convergence.

### 2.5.5 Stochastic Gradient Descent

**Stochastic gradient descent (SGD)**: update the parameters based on the gradient for a single training example

1- Choose  $i$  uniformly at random

$$2- \theta \leftarrow \theta - \alpha \frac{\partial \mathcal{L}^{(i)}}{\partial \theta}$$

- Cost of each SGD update is independent of  $N$ !
- SGD can make significant progress before even seeing all the data!
- Mathematical justification: if you sample a training example uniformly at random, the stochastic gradient is an **unbiased estimate** of the batch gradient:


$$\mathbb{E} \left[ \frac{\partial \mathcal{L}^{(i)}}{\partial \theta} \right] = \frac{1}{N} \sum_{i=1}^N \frac{\partial \mathcal{L}^{(i)}}{\partial \theta} = \frac{\partial \mathcal{J}}{\partial \theta}$$

### 2.5.6 Mini-batch Stochastic Gradient Descent

- **Problems** with using single training example to estimate gradient:
  - Variance in the estimate may be high
  - We can't exploit efficient vectorized operations
- **Compromise approach**:
  - Compute the gradients on a randomly chosen medium-sized set of training examples  $\mathcal{M} \subset \{1, \dots, N\}$  called a **mini-batch**.
- Stochastic gradients computed on larger mini-batches have smaller variance.
- The mini-batch size  $|\mathcal{M}|$  is a hyperparameter that needs to be set.
  - Too large: requires more compute: e.g., it takes more memory to store the activations, and longer to compute each gradient update
  - Too small: can't exploit vectorization, has high variance
  - A reasonable value might be  $|\mathcal{M}| = 100$ .

### 2.5.7 Comparison

- Batch gradient descent moves directly downhill (locally speaking).
- SGD takes steps in a noisy direction, but moves downhill on average.



- ▲ **Batch Gradient Descent**, 全批量梯度下降, 是最原始的形式, 它是指在每一次迭代时使用所有样本来进行梯度的更新。优点是全局最优解, 易于并行实现; 缺点是当样本数目很大时, 训练过程会很慢。
- ▲ **Stochastic Gradient Descent**, 随机梯度下降, 是指在每一次迭代时使用一个样本来进行参数的更新。优点是训

练速度快；缺点是准确度下降，并且可能无法收敛或者在最小值附近震荡。

▲ Mini-Batch Gradient Descent, 小批量梯度下降, 是对上述两种方法的一个折中办法。它是指在每一次迭代时使用一部分样本来进行参数的更新。这种方法兼顾了计算速度和准确度。

### 3 Linear Classifiers, Logistic Regression, Multiclass Classification

#### 3.1 Binary linear classification

- Classification: given a  $D$ -dimensional input  $\mathbf{x} \in \mathbb{R}^D$  predict a discrete-valued target
- Binary: predict a binary target  $t \in \{0,1\}$ 
  - Training examples with  $t = 1$  are called positive examples, and training examples with  $t = 0$  are called negative examples.
  - $t \in \{0,1\}$  or  $t \in \{-1, +1\}$  is for computational convenience.
- Linear: model prediction  $y$  is a linear function of  $\mathbf{x}$ , followed by a threshold  $r$ :

$$z = \mathbf{w}^T \mathbf{x} + b$$

$$y = \begin{cases} 1 & \text{if } z \geq r \\ 0 & \text{if } z < r \end{cases}$$

- **Eliminating the threshold:** We can assume without loss of generality (WLOG) that the threshold  $r = 0$

$$\mathbf{w}^T \mathbf{x} + b \geq r \iff \underbrace{\mathbf{w}^T \mathbf{x} + b - r}_{\triangleq w_0} \geq 0$$



- **Eliminating the bias:** Add a dummy feature  $x_0$  which always takes the value 1. the weight  $w_0 = b$  is equivalent to a bias (same as linear regression)
- Simplified model: receive input  $\mathbf{x} \in \mathbb{R}^{D+1}$  with  $x_0 = 1$

$$z = \mathbf{w}^T \mathbf{x}$$

- **Example:**

$x_0$	$x_1$	$t$
1	0	1
1	1	0

- Suppose this is our training set, with the dummy feature  $x_0$  included.
- Which conditions on  $w_0, w_1$  guarantee perfect classification?
  - When  $x_1 = 0$ , need:  $z = w_0 x_0 + w_1 x_1 \geq 0 \iff w_0 \geq 0$
  - When  $x_1 = 1$ , need:  $z = w_0 x_0 + w_1 x_1 < 0 \iff w_0 + w_1 < 0$
  - Example solution:  $w_0 = 1, w_1 = -2$





$$\begin{aligned} w_0 &\geq 0 \\ w_0 + w_1 &< 0 \end{aligned}$$

- Training examples are points
- Weights (hypotheses)  $\mathbf{w}$  can be represented by **half-spaces**:  $H^+ = \{\mathbf{x}: \mathbf{w}^T \mathbf{x} \geq 0\}$ ,  $H^- = \{\mathbf{x}: \mathbf{w}^T \mathbf{x} < 0\}$
- The boundary is the **decision boundary**:  $\{\mathbf{x}: \mathbf{w}^T \mathbf{x} = 0\}$
- If the training examples can be perfectly separated by a linear decision rule, we say **data is linearly separable**.
- Weights (hypotheses)  $\mathbf{w}$  are points
- Each training example  $\mathbf{x}$  specifies a half-space  $\mathbf{w}$  must lie in to be correctly classified:  $\mathbf{w}^T \mathbf{x} \geq 0$  if  $t = 1$ 
  - $x_0 = 1, x_1 = 0, t = 1 \implies (w_0, w_1) \in \mathbf{w}: w_0 \geq 0$
  - $x_0 = 1, x_1 = 1, t = 0 \implies (w_0, w_1) \in \mathbf{w}: w_0 + w_1 < 0$
- The region satisfying all the constraints is the feasible region; if this region is nonempty, the problem is



feasible, otherwise it is infeasible.



- Slice for  $x_0 = 1$  and
- Example sol:  $w_0 = -1.5, w_1 = 1, w_2 = 1$
- Decision boundary:
 
$$w_0x_0 + w_1x_1 + w_2x_2 = 0$$

$$\implies -1.5 + x_1 + x_2 = 0$$
- Slice for  $w_0 = -1.5$  for the constraints
- $w_0 < 0$
- $w_0 + w_2 < 0$
- $w_0 + w_1 < 0$
- $w_0 + w_1 + w_2 \geq 0$

### 3.2 Towards Logistic Regression

Define loss function then try to minimize the resulting cost function

#### Attempt 1: 0-1 loss

$$\mathcal{L}_{0-1}(y, t) = \begin{cases} 0 & \text{if } y = t \\ 1 & \text{if } y \neq t \end{cases}$$

$$= \mathbb{I}[y \neq t]$$


- Usually, the cost  $\mathcal{J}$  is the averaged loss over training examples; for 0-1 loss, this is the misclassification rate:

$$\mathcal{J} = \frac{1}{N} \sum_{i=1}^N \mathbb{I}[y^{(i)} \neq t^{(i)}]$$

- Minimum of a function will be at its critical points, use **Chain rule** to find the critical point of 0-1 loss

$$\frac{\partial \mathcal{L}_{0-1}}{\partial w_j} = \frac{\partial \mathcal{L}_{0-1}}{\partial z} \frac{\partial z}{\partial w_j}$$

- $\partial \mathcal{L}_{0-1} / \partial z$  is zero everywhere it's defined:




- $\partial \mathcal{L}_{0-1} / \partial w_j = 0$  means that changing the weights by a very small amount probably has no effect on the loss.
- Almost any point has 0 gradient!

#### Attempt 2: Linear Regression

$$z = \mathbf{w}^T \mathbf{x}$$

$$\mathcal{L}_{SE}(z, t) = \frac{1}{2} (z - t)^2$$

- Doesn't matter that the targets are actually binary. Treat them as continuous values.
- For this loss function, it makes sense to make final predictions by thresholding  $z$  at  $1/2$



- The loss function hates when you make correct predictions with high confidence!

- It  $t = 1$ , it's more unhappy about  $z = 10$  than  $z = 0$ .

### Attempt 3: Logistic Activation Function

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$


- $\sigma^{-1}(y) = \log(y/(1 - y))$  is called the **logit**.
- A linear model with a logistic nonlinearity is known as **log-linear**:

$$z = \mathbf{w}^T \mathbf{x}$$

$$y = \sigma(z)$$

$$\mathcal{L}_{SE}(y, t) = \frac{1}{2} (y - t)^2$$


- Used in this way,  $\sigma$  is called an **activation function**.



(plot of  $\mathcal{L}_{SE}$  as a function of  $z$ , assuming  $t = 1$ )

$$\frac{\partial \mathcal{L}}{\partial w_j} = \frac{\partial \mathcal{L}}{\partial z} \frac{\partial z}{\partial w_j}$$

- For  $z \ll 0$ , we have  $\sigma(z) \approx 0$ .
- $\partial \mathcal{L} / \partial z \approx 0$  (check)  $\rightarrow \partial \mathcal{L} / \partial w_j \approx 0 \rightarrow$  derivative w.r.t.  $w_j$  is small  $\rightarrow w_j$  is like a critical point.
- If the prediction is really wrong, you should be far from a critical point (which is your candidate solution).
- Because  $y \in [0, 1]$ , we can interpret it as the estimated probability that  $t = 1$ . If  $t = 0$ , then we want to heavily penalize  $y \approx 1$ .
- Cross-entropy loss (aka log loss) captures this intuition:



$$\begin{aligned} \mathcal{L}_{CE}(y, t) &= \begin{cases} -\log y & \text{if } t = 1 \\ -\log(1 - y) & \text{if } t = 0 \end{cases} \\ &= -t \log y - (1 - t) \log(1 - y) \end{aligned}$$

- The logistic loss is a **convex function** in  $\mathbf{w}$ , so let's consider the **gradient descent** method.
  - Recall: we initialize the weights to something reasonable and repeatedly adjust them in the direction of steepest descent.
  - A standard initialization is  $\mathbf{w} = 0$ .

$$\mathcal{L}_{CE}(y, t) = -t \log y - (1 - t) \log(1 - y)$$

$$y = 1 / (1 + e^{-z}) \text{ and } z = \mathbf{w}^T \mathbf{x}$$

$$\begin{aligned} \frac{\partial \mathcal{L}_{CE}}{\partial w_j} &= \frac{\partial \mathcal{L}_{CE}}{\partial y} \cdot \frac{\partial y}{\partial z} \cdot \frac{\partial z}{\partial w_j} = \left( -\frac{t}{y} + \frac{1-t}{1-y} \right) \cdot y(1-y) \cdot x_j \\ &= (y - t) x_j \end{aligned}$$

Gradient descent (coordinate-wise) update to find the weights of logistic regression:

$$\begin{aligned}
 w_j &\leftarrow w_j - \alpha \frac{\partial \mathcal{J}}{\partial w_j} \\
 &= w_j - \frac{\alpha}{N} \sum_{i=1}^N (y^{(i)} - t^{(i)}) x_j^{(i)}
 \end{aligned}$$

Gradient descent updates for Linear regression and Logistic regression (both examples of generalized linear models):

$$\mathbf{w} \leftarrow \mathbf{w} - \frac{\alpha}{N} \sum_{i=1}^N (y^{(i)} - t^{(i)}) \mathbf{x}^{(i)}$$

### 3.3 Multiclass Classification and Softmax Regression

#### 3.3.1 Multiclass Classification

- Classification tasks with more than two categories
- Targets form a discrete set  $\{1, \dots, K\}$ .
- It's often more convenient to represent them as **one-hot vectors**, or a **one-of-K encoding**:

$$\mathbf{t} = (0, \dots, 0, 1, 0, \dots, 0) \in \mathbb{R}^K$$

$\underbrace{\hspace{10em}}$   
 entry  $k$  is 1

- We can start with a linear function of the inputs.

$$z_k = \sum_{j=1}^D w_{kj} x_j + b_k \quad \text{for } k = 1, 2, \dots, K$$

- Now there are  $D$  input dimensions and  $K$  output dimensions, so we need  $K \times D$  weights, which we arrange as a weight matrix  $\mathbf{W}$ .
- Also, we have a  $K$ -dimensional vector  $\mathbf{b}$  of biases. Then eliminate the bias  $\mathbf{b}$  by taking  $\mathbf{W} \in \mathbb{R}^{K \times (D+1)}$  and adding a dummy variable  $x_0 = 1$ .

$$\mathbf{z} = \mathbf{W}\mathbf{x} + \mathbf{b} \quad \text{or with dummy } x_0 = 1 \quad \mathbf{z} = \mathbf{W}\mathbf{x}$$

- We can interpret the magnitude of  $z_k$  as a measure of how much the model prefers  $k$  as its prediction to turn this linear prediction into a one-hot prediction.

$$y_i = \begin{cases} 1 & i = \arg \max_k z_k \\ 0 & \text{otherwise} \end{cases}$$

#### 3.3.2 Softmax Regression

- We need to soften our predictions for the sake of optimization.
- We want soft predictions that are like probabilities, i.e.,  $0 \leq y_k \leq 1$  and  $\sum_k y_k = 1$ .
- A natural activation function to use is the softmax function, a multivariable generalization of the logistic function:

$$y_k = \text{softmax}(z_1, \dots, z_K)_k = \frac{e^{z_k}}{\sum_{k'} e^{z_{k'}}}$$

- Outputs can be interpreted as probabilities (positive and sum to 1)
- If  $z_k$  is much larger than the others, then  $\text{softmax}(\mathbf{z})_k \approx 1$  and it behaves like argmax.

- If a model outputs a vector of class probabilities, we can use cross-entropy as the loss function:

$$\begin{aligned}
 \mathcal{L}_{\text{CE}}(\mathbf{y}, \mathbf{t}) &= - \sum_{k=1}^K t_k \log y_k \\
 &= -\mathbf{t}^T (\log \mathbf{y})
 \end{aligned}$$


- Just like with logistic regression, we typically combine the softmax and cross-entropy into a **softmax-cross-entropy function**.
- Softmax regression (with dummy  $x_0 = 1$ ):

$$\begin{aligned}
 \mathbf{z} &= \mathbf{W}\mathbf{x} \\
 \mathbf{y} &= \text{softmax}(\mathbf{z}) \\
 \mathcal{L}_{\text{CE}} &= -\mathbf{t}^T (\log \mathbf{y})
 \end{aligned}$$

- Gradient descent updates can be derived for each row of  $\mathbf{W}$ :

$$\frac{\partial \mathcal{L}_{\text{CE}}}{\partial \mathbf{w}_k} = \frac{\partial \mathcal{L}_{\text{CE}}}{\partial z_k} \cdot \frac{\partial z_k}{\partial \mathbf{w}_k} = (y_k - t_k) \cdot \mathbf{x}$$
$$\mathbf{w}_k \leftarrow \mathbf{w}_k - \alpha \frac{1}{N} \sum_{i=1}^N (y_k^{(i)} - t_k^{(i)}) \mathbf{x}^{(i)}$$

- Similar to linear/logistic reg (no coincidence) (verify the update)
- Sometimes we can overcome the nonlinear problem with feature maps:



## 4 L

### 4.1 B